

Classifying Eye State from EEG Signals via k -Nearest Neighbours: A Cross-Validation Honesty Study

Ivan Del Rio Anish Kondamadugula

COGS 109 — Modelling and Data Analysis (Mukamel), Spring 2026
University of California, San Diego

June 10, 2026

Abstract

We pick a single data-analysis approach from the COGS 109 study guide—**classification**—and a single model within that approach— **k -nearest neighbours (KNN)**—and we perform model selection by sweeping the hyperparameter k over a log-spaced grid $\{1, 3, 5, 7, 11, 15, 21, 31, 51, 75, 99, 151, 201\}$ to find the value that minimises 5-fold cross-validation (CV) error on the UCI #264 EEG Eye State dataset (Roesler, 2013). We then run that *same* model-selection procedure under three different cross-validation schemes: shuffled, naive blocked, and stratified blocked. The contribution is *methodological*: we document how the choice of CV scheme changes the selected hyperparameter and the reported accuracy on an autocorrelated EEG recording. All three schemes select $k = 1$, but the reported accuracies differ dramatically: $97.3\% \pm 0.4\%$ under shuffled CV, $50.0\% \pm 10.7\%$ under naive blocked CV, and $77.8\% \pm 2.7\%$ under honest stratified blocked CV. The 19.5 *pp* gap between the shuffled and stratified-blocked estimates is leakage attributable to scheme choice—not to model quality, since the identical KNN model is evaluated in all three settings on the same training data. A sanity-check appendix runs the same idea on three alternative classifiers (LDA, PCA→LDA, PCR-as-classifier) and finds leakage gaps of only 2–6 *pp*, consistent with KNN being uniquely vulnerable to lag-1 autocorrelation because it predicts from per-sample similarity.

1 Introduction

1.1 Problem and motivation

Electroencephalography (EEG) eye-state classification—deciding from raw scalp voltages whether a subject’s eyes are open or closed—is a building block of brain–computer interfaces, drowsiness detection, and attention-monitoring systems. The UCI #264 EEG Eye State dataset is one of the most widely used benchmarks for this task. The dataset’s originators report a best error of 2.7% ($\approx 97.3\%$ accuracy) using the instance-based KStar classifier under Weka’s default (shuffled) 10-fold cross-validation, while non-instance-based methods score far worse (Roesler & Suendermann, 2013); online tutorials and student write-ups likewise routinely report k -nearest-neighbour accuracies in the 95–98% range under standard shuffled k -fold cross-validation.

We argue in this report that those numbers reflect a *methodological artefact* rather than a strong classifier. The dataset is a single ~ 117 -second continuous recording from one subject; its label changes only 23 times across 14,980 samples; adjacent voltage samples are nearly identical (per-channel lag-1 autocorrelation ≈ 0.97); and the label is piecewise-constant over runs that are sometimes thousands of samples long (lag-1 *label* autocorrelation ≈ 0.997). A KNN classifier with $k = 1$ evaluated under shuffled k -fold CV is essentially asking “what was the eye state of the sample 8 ms before or after this one?”—a question to which it can almost always answer correctly. That is a property of the cross-validation rule, not of the classifier.




Figure 1: Class balance of the UCI #264 EEG eye-state dataset. The majority class (eyes open) sets the random-guess baseline at 0.5512.

The question this report addresses is therefore methodological rather than predictive:

Given an autocorrelated single-subject EEG recording and a KNN classifier, how does the choice of cross-validation scheme change the result of an otherwise identical model-selection procedure?

This matters because cross-validation is the standard tool students are taught for honest accuracy estimation, yet on autocorrelated time-series data the default i.i.d. shuffled scheme silently inflates the estimate. Quantifying that inflation on a clean, well-known benchmark makes the failure mode concrete and reproducible.

1.2 Dataset and key characteristics

We use the UCI EEG Eye State dataset (Roesler, 2013; UCI Machine Learning Repository #264): a single uninterrupted recording from one subject wearing an Emotiv EPOC consumer-grade EEG headset.

- **Observations:** $n = 14,980$ samples, recorded at ~ 128 Hz over ~ 117 s.
- **Predictors:** $p = 14$ continuous voltage channels (AF3, F7, F3, FC5, T7, P7, O1, O2, P8, T8, FC6, F4, F8, AF4).
- **Response:** the binary label `eyeDetection` (0 = eyes open, 1 = eyes closed), annotated by the experimenter.
- **Class balance:** 55.12% open (8,257 samples) vs. 44.88% closed (6,723 samples), so the majority-class baseline—the accuracy floor any useful classifier must beat—is 0.5512 (Figure 1).
- **Noise sources:** consumer-grade hardware prone to electrode drift, motion and muscle artefacts, and occasional headset disconnections (visible as a handful of extreme-voltage outliers).

A purely structural feature of this dataset is the heart of the study. The labels change rarely: the recording contains only **24 contiguous label runs** (the subject opens or closes their eyes only 23 times in ~ 117 s), with run lengths ranging from under a hundred to several thousand consecutive samples (Figure 2). Consequently the lag-1 autocorrelation of a single voltage channel is high (≈ 0.97 on average, ranging 0.91–0.99 across the 14 channels after outlier removal)—adjacent samples are nearly identical—and the label, being piecewise-constant over long runs, is even more persistent: its autocorrelation is ≈ 0.997 at lag 1, ≈ 0.83 at lag 50 (≈ 0.4 s), crossing zero only at very long lags (Figure 3).

The consequence is direct: *any* CV protocol that splits samples uniformly at random leaves each test sample with a near-duplicate in its own training set, so a model that predicts from per-sample




Figure 2: Raw 14-channel EEG voltages (top) with the eye-state label strip (bottom; blue = eyes open, 0; red = eyes closed, 1). Channels are continuous and smooth on the millisecond scale; the label is piecewise-constant over multi-second runs.




Figure 3: Autocorrelation of the binary `eyeDetection` label across lags 1–1000. The label is essentially identical to itself for the first ~ 100 samples (~ 0.8 s) and only decorrelates at very long lags. This is the structural reason shuffled CV leaks: a sample’s nearest neighbour in time is almost always in its own class.

similarity (KNN at $k = 1$) will look spectacular while merely memorising—not learning a generalisable voltage-to-label rule.

1.3 Hypothesis

Hypothesis. A KNN classifier trained on raw 14-channel EEG voltages can predict per-sample eye state above the 55.12% majority-class baseline, but accuracy estimates obtained under shuffled k -fold cross-validation will substantially exceed honest out-of-time estimates because of temporal leakage between adjacent samples.

The hypothesis has two parts: (i) the honest accuracy beats the majority baseline, and (ii) the shuffled estimate is inflated relative to a leakage-resistant estimate. We test both and report what we find.

2 Methods

2.1 Approach: classification (prediction)

The COGS 109 study guide (Mukamel, 2026) enumerates several broad approaches (linear, polynomial and spline regression, PCR, LDA, KNN, K -means, hierarchical clustering, PCA). The target variable here is the binary `eyeDetection` label, so **classification** is the natural approach: we want a model that

maps a 14-dimensional voltage vector to one of two discrete labels. Regression approaches do not target a discrete label without artificial thresholding, and clustering approaches ignore the label entirely. We are focused on **prediction** (out-of-sample accuracy of the eye-state label), not on inference about which channel “causes” an eye state; the methodological question is specifically about how honestly we can *estimate* that predictive accuracy.

2.2 Model: k -nearest neighbours

Within classification we chose KNN as our model (implemented with scikit-learn; Pedregosa et al., 2011). KNN assigns a test sample the majority label among its k nearest training samples under Euclidean distance in the z -scored 14-D channel space. We chose it for three reasons. First, it is the simplest non-parametric classifier in the palette: it makes no distributional assumption and has exactly one hyperparameter (k) controlling the bias–variance trade-off (James et al., 2021), making it ideal for a clean single-axis hyperparameter sweep. Second—and most importantly for this study—because KNN predicts from per-sample similarity to the training set, it is the model in the palette that *most directly* exploits the lag-1 autocorrelation of EEG data; if we want to study how CV scheme choice changes a result, picking the model most sensitive to that choice gives the cleanest empirical signal. Third, KNN at small k reproduces the $\sim 97\%$ literature numbers on this dataset under shuffled CV, so the leakage gap we measure is a like-for-like comparison between an honest and a leaky estimate of the same baseline.

2.3 Preprocessing

1. **Outlier removal.** We drop the four samples with at least one channel exceeding 4 standard deviations above the channel mean ($z > 4$). This is standard EEG hygiene for headset disconnections and artefacts; an ablation (`tables/02_outlier_ablation.csv`) shows it does not materially change downstream accuracy but stabilises per-channel summary statistics. Cleaning leaves 14,976 samples.
2. **Chronological split.** We split the cleaned recording chronologically into 80% train ($n = 11,917$) and 20% test ($n = 2,995$), inserting a **64-sample seam gap** (0.5 s at 128 Hz) between train and test so no train sample is within half a second of a test sample. This is the minimum cross-partition isolation warranted by the autocorrelation profile in Figure 3.
3. **z -scoring (train statistics only).** We z -score each channel using the *training* mean and standard deviation only, and reuse those parameters for the test set. z -scoring before the split would itself leak (the test set would shape the channel scale), so we defer it to after the split.

2.4 Model-selection procedure: the k -sweep

The hyperparameter to tune for KNN is k , the number of neighbours that vote. We sweep

$$k \in \{1, 3, 5, 7, 11, 15, 21, 31, 51, 75, 99, 151, 201\},$$

a log-spaced, odd-only grid (odd avoids ties on a binary label) that starts at $k = 1$ (the most leakage-prone setting) and extends to $k = 201$ (a local average that should damp out lag-1 similarity if a real classification signal exists). For each k and each CV scheme we run 5-fold cross-validation on the z -scored training partition and aggregate per-fold accuracy as mean $\pm$ standard deviation. The **selected** k under a scheme is the argmax of mean CV accuracy. Crucially, the selection *rule* (argmax accuracy) is identical across schemes—only the fold-assignment rule changes.

Two caveats on how we report this. First, the accuracy at the argmax-selected k slightly overestimates true performance, because the maximum over a noisy grid is upward-biased (Cawley & Talbot, 2010); here the effect is negligible, as the stratified-blocked curve is nearly flat near the optimum ($k=1: 0.7778$, $k=3: 0.7777$, $k=5: 0.7770$). Second, the $\pm$ we report is the standard deviation of accuracy across the five folds from a single seed; because the folds share training data this is an informal measure of dispersion, not a calibrated confidence interval (Bengio & Grandvalet, 2004).

2.5 Cross-validation schemes (how the splitting actually works)

We implement all three schemes ourselves rather than delegating to a library, so we can state exactly how each one partitions the $N = 11,917$ training samples into 5 folds. In every scheme each sample appears in exactly one test fold, and a fold’s training set is the complement of its test set.

Scheme A — shuffled 5-fold (leaky baseline). We generate an index array $[0, 1, \dots, N - 1]$ and apply a single seed-fixed Fisher–Yates permutation (NumPy `default_rng(42)`). The shuffled indices are cut into 5 contiguous equal blocks (the last block absorbs the remainder); block i is the test set of fold i and the other four blocks are its training set. Because the permutation is uniform, each test fold is a random 20% subset spread across the whole recording—so the immediate temporal neighbours of almost every test sample land in the training folds. This is the standard “`KFold(shuffle=True)`” protocol, and the leakiest possible scheme on a continuous time series.

Scheme B — naive blocked 5-fold. We take the time-ordered index array *without* shuffling and cut it into 5 contiguous chunks. Fold i ’s test set is chunk i (a single contiguous span of time); its training set is the remaining time. This eliminates lag- ≤ 1 leakage but introduces a new problem: because the label runs are long and unevenly distributed, individual chunks can be near class-pure (e.g. one chunk at $\sim 20\%$ class-0 vs. the overall $\sim 54\%$). Per-fold accuracy variance therefore explodes and the mean can fall *below* the majority baseline. We include it to make the failure mode visible—it is *not* a scheme we recommend.

Scheme C — stratified blocked 5-fold (honest). This is the protocol we recommend. We cut the time-ordered training partition into $M = 100$ short contiguous segments of ~ 119 samples (~ 0.93 s) each, preserving within-segment temporal continuity. We label each segment by its class-1 proportion, sort the segments by that proportion, and walk the sorted list in chunks of 5; within each chunk the five segments are assigned to the five folds by a seed-fixed random permutation. The effect is that each fold receives one segment from each quintile of the class-1-proportion distribution, so every fold’s class balance is within ~ 3 pp of the overall $\sim 54\%$ class-0 balance. This keeps the temporal-isolation guarantee of blocked CV (neighbours stay together inside a segment) *and* restores the class-balance guarantee of stratified CV. One residual leak remains: we do not insert a seam gap *between* adjacent segments, so the two samples straddling each of the ~ 99 segment boundaries can land in different folds. This affects under 1% of training samples ($99/11,917$), so it can inflate the stratified-blocked accuracy by at most a point or two—a caveat we revisit when reconciling the CV and holdout estimates (§3.5). Figure 4 shows the chronological training partition with the naive-blocked 5-fold boundaries and the eye-state label strip, making the long, unevenly distributed label runs—the reason naive blocking fails and stratification is needed—visible at a glance.

2.6 Comparison models (sanity checks)

To test whether the leakage gap is KNN-specific or a generic property of this dataset, we run the same model-selection idea on three other classifiers from the palette. **LDA** has no tunable hyperparameter—one accuracy per scheme. **PCA**→**LDA** sweeps $n_{\text{comp}} \in \{1, 2, 3, 5, 7, 10\}$ and selects the n maximising mean stratified-blocked CV accuracy. **PCR-as-classifier** (ordinary least squares on the first n principal components, thresholded at 0.5) sweeps the same grid. These are supporting evidence, not the primary result.

2.7 Metrics

We report **accuracy** = $(\text{TP} + \text{TN})/N$ for both CV and holdout, and—because the chronological holdout is heavily class-imbalanced (90.6% class-0, as the final segment of the recording is a long eyes-open run)—we also report **sensitivity** (recall for closed = 1), **specificity** (recall for open = 0), and **balanced accuracy** (the mean of the two) on the holdout, since raw accuracy on a 91/9 split is misleading. We distinguish two kinds of estimate throughout: the *within-recording* CV accuracy (resampled from the




Figure 4: The chronological training partition (z -scored 14 channels, offset for readability) with the four naive-blocked 5-fold boundaries (red vertical lines) and the eye-state label strip (blue = eyes open, 0; red = eyes closed, 1). The long, unevenly distributed label runs mean the five contiguous naive-blocked folds are far from class-balanced, motivating the stratified-blocked scheme (Scheme C), which instead redistributes 100 short contiguous segments across folds.

Table 1: Per-scheme selected k and accuracy from the KNN k -sweep. All three schemes select $k = 1$; the reported accuracies span 47 percentage points.

CV scheme	Selected k	Mean accuracy $\pm$ SD
Shuffled (leaky baseline)	1	0.9728 ± 0.0037 ($\approx 97.3\%$)
Naive blocked	1	0.5004 ± 0.1068 ($\approx 50.0\%$)
Stratified blocked (honest)	1	0.7778 ± 0.0274 ($\approx 77.8\%$)

training partition) and the *out-of-time* holdout accuracy (a genuinely future, unseen segment); §3.5 shows they disagree, and we discuss why.

3 Results

3.1 Model selection: KNN under three CV schemes

Figure 5 is the headline figure. It plots KNN mean accuracy $\pm$ standard deviation as a function of k (log scale) under all three CV schemes on the same axes.

Three observations are immediate. The shuffled curve sits near 0.97 across all k , decaying roughly linearly on the log axis from 0.973 at $k = 1$ to ~ 0.84 at $k = 201$ —“nearly perfect” everywhere. The naive blocked curve sits near 0.50 across all k , at or below the majority baseline, with very large per-fold standard deviation (~ 10 –13 pp); this is fold construction failing, not the model failing (§4.3). The stratified blocked curve peaks at 0.778 at $k = 1$ and decays gracefully to ~ 0.71 by $k = 201$, with a much smaller standard deviation (~ 2.7 –4.0 pp).

Applying the argmax-accuracy rule to each scheme yields Table 1.

The remarkable result is that **all three schemes select the same hyperparameter** ($k = 1$): the model-selection procedure is robust, but the *answer* it reports is not. The selected accuracy ranges from 50.0% (naive blocked) to 97.3% (shuffled)—a 47 pp spread—despite the procedure, the model, the data, and the selected k being identical.

3.2 Complexity, flexibility, and overfitting

For KNN, model flexibility is inversely related to k : $k = 1$ is the most flexible (a jagged, high-variance decision boundary that can fit any training point), and large k is more rigid (a smooth, high-bias boundary). The shape of the curves in Figure 5 is itself diagnostic of overfitting-via-leakage. Under




Figure 5: KNN model selection under three CV schemes (5-fold each; error bars are ± 1 SD across folds). Orange squares: shuffled CV (leaky baseline). Blue circles: naive blocked CV. Green diamonds: stratified blocked CV (honest). The horizontal dashed line is the majority-class baseline (0.5512). The open marker on each curve highlights the argmax-accuracy k and the corresponding accuracy.

honest stratified blocked CV, accuracy *decreases* monotonically as k grows from 1 to 201 ($0.778 \rightarrow 0.710$): the flexible model generalises best, because the signal it exploits (short-range temporal smoothness within a segment) is real and local. Under *shuffled* CV the curve is pathologically flat-and-high (~ 0.97 everywhere): the model memorises the temporal twin of each test sample regardless of k . That flat-near-1 shape is the fingerprint of leakage-driven overfitting—the estimator looks like it cannot overfit because the “test” samples are near-duplicates of training samples.

3.3 Leakage-gap analysis

The methodological centrepiece is the **leakage gap**: the difference between an estimator’s leaky shuffled-CV accuracy and its honest stratified-blocked-CV accuracy at the same selected hyperparameter. For KNN at $k = 1$,

$$\text{leakage gap} = 0.9728 - 0.7778 = 0.1950 \quad (\approx 19.5 \text{ pp}).$$

We attribute this gap to leakage of adjacent samples across the train/test fold boundary, not to model quality: the model is 1-NN in both cases and the training partition is identical—only the fold-assignment rule changes. Under shuffled CV the temporal neighbour of every test sample lands in a training fold with $\sim 80\%$ probability and carries the same label with $\sim 99.7\%$ probability (the lag-1 *label* autocorrelation); under stratified blocked CV that neighbour is reserved into the test sample’s own segment in all but the $<1\%$ of cases that straddle a segment boundary. The honest within-recording accuracy at $k = 1$, $77.8\% \pm 2.7\%$, is ~ 23 pp above the 55.12% baseline, so the shuffled estimate is clearly inflated (hypothesis part (ii)). Whether this within-recording signal *generalises out-of-time* is a separate question that the holdout answers less favourably (§3.5).

3.4 Supporting evidence: alternative classifiers

To rule out a generic explanation, we ran the same idea on three other classifiers (Table 2, Figure 6).

The pattern is exactly what we expect if the gap is caused by per-sample autocorrelation. LDA’s decision rule is a single global linear projection of the 14-D space onto a 1-D axis; it cannot exploit per-sample similarity, so its gap is small (5.6 pp). PCA→LDA and PCR both project to a low-dimensional

Table 2: Per-model leakage gaps (shuffled CV minus stratified-blocked CV) at the selected hyperparameter. KNN’s gap is the outlier: $\sim 4\text{--}8\times$ larger than any other classifier in the palette.

Model	Selected hyperparam.	Shuffled CV	Strat. blocked CV	Leakage gap
LDA	— (no sweep)	0.6471 ± 0.0063	0.5911 ± 0.0474	$+5.6\text{ pp}$
PCA→LDA	$n = 3$	0.5574 ± 0.0147	0.5339 ± 0.0438	$+2.4\text{ pp}$
PCR (classifier)	$n = 2$	0.5528 ± 0.0087	0.5250 ± 0.0176	$+2.8\text{ pp}$
KNN	$k = 1$	0.9728 ± 0.0037	0.7778 ± 0.0274	$+19.5\text{ pp}$




Figure 6: Three-way 5-fold CV accuracy across the four palette classifiers. Orange: shuffled; blue: naive blocked; green: stratified blocked. The KNN bar group has by far the largest shuffled-to-stratified-blocked gap, consistent with KNN being uniquely vulnerable to lag-1 autocorrelation.

subspace before deciding, destroying most of the high-frequency sample-to-sample similarity KNN relies on; their gaps are smaller still (2.4 and 2.8 *pp*). KNN at $k = 1$ is the only classifier whose prediction depends on the identity of the single nearest training sample—and its gap is $4\text{--}8\times$ larger than any other.

3.5 Model estimation: final parameters and accuracy

Final model and “parameters”. The selected model is KNN with $k = 1$ on the z -scored 14-channel features. KNN is non-parametric, so the “fitted parameters” are not coefficients but (i) the selected hyperparameter $k = 1$ and (ii) the stored training set together with the per-channel z -scoring constants (μ_c, σ_c) . Following the standard convention of refitting the final model on all available data once the hyperparameter is fixed (James et al., 2021), the deployable model recomputes (μ_c, σ_c) and sets the 1-NN store on the *entire* cleaned dataset ($n = 14,976$); cross-validation is used only for selection, not for the final fit. For a concrete parametric contrast, the LDA comparison model *does* have estimable coefficients ([figures/10_lda_coefficients.png](#)): the largest-magnitude discriminant weights fall on the temporal (T7, T8) and frontal (F8, F7) channels, with parietal P7 also prominent, while the occipital channels O1/O2 are small-to-negligible. We therefore do *not* attribute the discriminant to posterior alpha rhythms; large temporal and frontal weights are more consistent with eye-movement and muscle (EMG) artefact and with the per-sensor amplitude-range differences that Roesler & Suendermann (2013) identify as discriminative on this dataset. A linear discriminant on raw per-sample voltages cannot in any case isolate band-limited alpha power.

Final prediction accuracy. The headline estimate is the honest *within-recording* cross-validated accuracy of the selected model ($k = 1$, stratified blocked CV): **$77.8\% \pm 2.7\%$** , versus the 55.12% majority baseline.




Figure 7: Per-model confusion matrices on the 20% chronological holdout test set (rows: true; columns: predicted). The holdout is dominated by eyes-open samples ($\sim 91\%$ class-0), so accuracy here is mostly driven by specificity. The number in each panel is that model’s holdout accuracy.

Reconciling the holdout: out-of-time generalisation is weaker. The 20% chronological holdout is the single most leakage-resistant test in the study—it is genuinely future data, separated from training by the 64-sample seam gap—so we confront it directly rather than discarding it. On the holdout, KNN ($k = 1$) scores only 51.3% accuracy with sensitivity 0.45 and specificity 0.52, i.e. a balanced accuracy of ≈ 0.49 : essentially chance, and far below the holdout’s own 90.6% majority baseline (the final segment is a long eyes-open run, so the holdout is 90.6% class-0 versus 46.5% in training; Figure 7). The within-recording CV cannot see this prior shift because every fold is rebalanced to the training distribution, whereas the holdout can. The honest reading is therefore split: the leakage result (part (ii)) is robust, but the claim that the model captures a *generalisable* eye-state signal (part (i)) holds only *within* the recording and is *not* corroborated out-of-time on this single ~ 117 s session.

4 Conclusions and Discussion

4.1 What we conclude about the hypothesis

Our hypothesis holds in part. Part (ii) is strongly supported: the shuffled-CV estimate (97.3%) is inflated by 19.5 *pp* relative to the honest stratified-blocked estimate, and that inflation is leakage attributable purely to the CV scheme—the model, the selected k , and the training data are identical across schemes (§3.3). Part (i) is supported only within the recording: the honest *within-recording* KNN accuracy (77.8%) beats the 55.12% baseline by ~ 23 *pp*, but the out-of-time holdout (balanced accuracy ≈ 0.49 , §3.5) does *not* corroborate a generalisable voltage-to-state rule. We report this candidly: on a single-subject, single-session recording, within-recording resampling and genuine out-of-time prediction disagree, and only the leakage finding is unambiguous.

Shuffled CV says we have a 97% classifier. Honest stratified blocked CV says we have a 78% classifier. The difference is leakage, not signal.

4.2 Why KNN is uniquely vulnerable

The mechanism is the one quantified in §3.3: because the per-channel lag-1 voltage autocorrelation is ≈ 0.97 , a test sample’s nearest neighbour is almost always its temporal twin, whose label matches with probability ≈ 0.997 . Shuffled CV leaves that twin in a training fold, so 1-NN exceeds 95% by memorisation; stratified blocked CV reserves it into the test sample’s own segment, dropping accuracy to 77.8%. Models that average over many samples (LDA, PCA→LDA, PCR) cannot exploit per-sample twins and show only 2–6 pp gaps.

4.3 Why naive blocked CV underperforms the baseline

The five contiguous chunks are not class-balanced: with only 23 label transitions in the recording, several chunks straddle few or zero transitions, and per-chunk class proportions stray far from the global 55/45 split. A classifier trained on four mostly-open chunks then predicts “open” on a held-out mostly-closed chunk, scoring *below* the majority baseline. This is a fold-stratification failure, not a model failure; stratifying contiguous segments (Scheme C) fixes it, which is why stratified blocked CV recovers ~ 28 pp over naive blocked.

4.4 Implications and future work

A non-trivial fraction of tutorials and student projects on this dataset report 95–98% KNN accuracy under shuffled k -fold. Our analysis suggests those numbers are a CV-protocol artefact: the honest within-recording accuracy is 77.8% and the extra ~ 20 pp is adjacent-sample leakage. Tellingly, even the dataset’s originators obtain their best results with instance-based learners (KStar/IB1, $\approx 97\%$) under shuffled 10-fold CV while non-instance-based methods score far worse (Roesler & Suendermann, 2013)—exactly the signature of leakage that instance-based models are uniquely able to exploit. This echoes broader cautions about i.i.d. CV on autocorrelated signals (Bergmeir & Benítez, 2012). The practical recommendation is to use blocked, class-stratified CV (or an explicit out-of-time holdout with a seam gap) whenever samples are temporally autocorrelated. Limitations of our study: it is a single subject on consumer hardware over ~ 117 s, and we use per-sample voltages with no windowed features. Future work should (i) add windowed features (rolling means/variances, lagged values), which would give KNN more meaningful neighbours and likely shrink the leakage gap; (ii) test cross-subject generalisation on a multi-subject corpus; and (iii) extend the k sweep to characterise the leakage curve more finely. The project is, above all, a worked example of how a methodologically sound model-selection procedure can still be misled by an upstream choice—the CV scheme—unrelated to the model itself.

References

1. Roesler, O. (2013). *EEG Eye State Data Set*. UCI Machine Learning Repository, id #264. <https://archive.ics.uci.edu/dataset/264/eeg+eye+state>
2. Roesler, O., & Suendermann, D. (2013). A First Step towards Eye State Prediction Using EEG. *Proc. AIHLS*.
3. Cawley, G. C., & Talbot, N. L. C. (2010). On over-fitting in model selection and subsequent selection bias in performance evaluation. *Journal of Machine Learning Research* 11:2079–2107.
4. Bengio, Y., & Grandvalet, Y. (2004). No unbiased estimator of the variance of k -fold cross-validation. *Journal of Machine Learning Research* 5:1089–1105.
5. Mukamel, R. (2026). *COGS 109 — Modelling and Data Analysis, Spring 2026 Study Guide*. UC San Diego.
6. James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An Introduction to Statistical Learning*, 2nd ed. Springer.
7. Bergmeir, C., & Benítez, J. M. (2012). On the use of cross-validation for time series predictor evaluation. *Information Sciences* 191:192–213.
8. Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research* 12:2825–2830.